

浙江大学

操作系统实验五

课程名称:	操作系统
作业名称:	RV64 缺页异常处理与 fork 机制
姓 名:	仇国智
学 号:	3220102181
电子邮箱:	3220102181@zju.edu.cn
联系电话:	13714176104
指导教师:	申文博

2024 年 11 月 24 日

目录

1	实验内容	1
2	实验原理	1
2.1	vm_area_struct 介绍	1
2.2	缺页异常 page fault	2
2.3	Demand paging	2
2.4	RISC-V Page Faults	3
2.5	处理 page fault 的方式	3
2.6	Fork 系统调用	4
3	实验具体过程与代码实现	4
3.1	缺页异常处理	4
3.1.1	实现虚拟内存管理功能	4
3.1.2	修改 tas_init	6
3.1.3	实现 page fault handler	8
3.2	实现 fork 系统调用	11
3.2.1	拷贝内核栈	11
3.2.2	创建子进程页表	11
3.2.3	处理进程返回逻辑	13
3.2.4	写时复制 COW	14
4	实验结果与分析	17
4.1	测试缺页处理	17
4.2	测试 fork	20
5	遇到的问题及解决方法	23
6	总结与心得	23
6.1	画图分析 make run TEST=FORK3 的进程 fork 过程, 并呈现出各个进程的 global_variable 应该从几开始输出, 再与你的输出进行对比验证.	23

1 实验内容

- 通过 `vm_area_struct` 数据结构实现对进程多区域虚拟内存的管理
- 在 Lab4 实现用户态程序的基础上, 添加缺页异常处理 `page fault handler`
- 为进程加入 `fork` 机制, 能够支持通过 `fork` 创建新的用户态进程

2 实验原理

2.1 `vm_area_struct` 介绍

在 Linux 系统中,`vm_area_struct` 是虚拟内存管理的基本单元,`vm_area_struct` 保存了有关连续虚拟内存区域 (简称 `vma`) 的信息.Linux 具体某一进程的虚拟内存区域映射关系可以通过 `procfs` 读取 `/proc/<pid>/maps` 的内容来获取.

在本次实验中, 我们需要实现对进程多区域虚拟内存的管理, 即通过 `vm_area_struct` 数据结构实现对进程多区域虚拟内存的管理. 在 Lab4 实现用户态程序的基础上, 添加缺页异常处理 `page fault handler`, 为进程加入 `fork` 机制, 能够支持通过 `fork` 创建新的用户态进程.

```
$ cat /proc/7884/maps
556f22759000-556f22786000 r--p 00000000 08:05 16515165          /usr/bin/bash
556f22786000-556f22837000 r-xp 0002d000 08:05 16515165          /usr/bin/bash
556f22837000-556f2286e000 r--p 000de000 08:05 16515165          /usr/bin/bash
556f2286e000-556f22872000 r--p 00114000 08:05 16515165          /usr/bin/bash
556f22872000-556f2287b000 rw-p 00118000 08:05 16515165          /usr/bin/bash
556f22fa5000-556f2312c000 rw-p 00000000 00:00 0               [heap]
7fb9edb0f000-7fb9edb12000 r--p 00000000 08:05 16517264          /usr/lib/x86_64-linux-gnu/libnss_files-2.31.so
7fb9edb12000-7fb9edb19000 r-xp 00003000 08:05 16517264          /usr/lib/x86_64-linux-gnu/libnss_files-2.31.so
...
7ffee5cdc000-7ffee5cfd000 rw-p 00000000 00:00 0               [stack]
7ffee5dce000-7ffee5dd1000 r--p 00000000 00:00 0               [vvar]
7ffee5dd1000-7ffee5dd2000 r-xp 00000000 00:00 0               [vdso]
fffffffffff600000-fffffffffff601000 --xp 00000000 00:00 0      [vsyscall]
```

从中我们可以读取如下一些有关该进程内虚拟内存映射的关键信息:

- `vm_start`:(第 1 列) 该段虚拟内存区域的开始地址
- `vm_end`:(第 2 列) 该段虚拟内存区域的结束地址
- `vm_flags`:(第 3 列) 该段虚拟内存区域的一组权限 (`rwX`) 标志,`vm_flags` 的具体取值定义可参考 Linux 源代码的 `linux/mm.h`
- `vm_pgoff`:(第 4 列) 虚拟内存映射区域在文件内的偏移量

- `vm_file`: (第 5/6/7 列) 分别表示: 映射文件所属设备号 / 以及指向关联文件结构的指针 / 以及文件名

我们注意到, 一段内存中的内容可能是由磁盘中的文件映射的. 如果这样的内存的 VMA 产生了缺页异常, 说明文件中对应的页不在操作系统的 buffer pool 中, 或者是由于 buffer pool 的调度策略被换出到磁盘上了. 这时候操作系统会用驱动读取硬盘上的内容, 放入 buffer pool, 然后修改当前进程的页表来让其能够用原来的地址访问文件内容, 而这一切对用户程序来说是完全透明的, 除了访问延迟.

除了跟文件建立联系以外, VMA 还可能是一块匿名 (anonymous) 的区域. 例如被标成 [stack] 的这一块区域, 并没有对应的文件.

其它保存在 `vm_area_struct` 中的信息还有:

- `vm_ops`: 该 `vm_area` 中的一组工作函数, 其中是一系列函数指针, 可以根据需要进行定制
- `vm_next/vm_prev`: 同一进程的所有虚拟内存区域由链表结构链接起来, 这是分别指向前后两个 `vm_area_struct` 结构体的指针

2.2 缺页异常 page fault

在一个启用了虚拟内存的系统上, 若正在运行的程序访问当前未由内存管理单元 (MMU) 映射到虚拟内存的页面, 或访问权限不足, 则会由计算机硬件引发的缺页异常 (page fault).

处理缺页异常通常是操作系统内核的一部分, 当处理缺页异常时, 操作系统将尝试使所需页面在物理内存中的位置变得可访问 (建立新的映射关系到虚拟内存). 而如果在非法访问内存的情况下, 即发现触发 page fault 的虚拟内存地址 (Bad Address) 不在当前进程的 `vm_area_struct` 链表中所定义的允许访问的虚拟内存地址范围内, 或访问位置的权限条件不满足时, 缺页异常处理将终止该程序的继续运行.

2.3 Demand paging

Demand paging 是一种虚拟内存管理技术, 它允许程序在需要时才将页面加载到内存中. 这样做的好处是, 仅加载执行进程所需的页面, 从而节省内存空间. 例如, 若一个页面从未被访问过, 那么它就不需要被放入内存中.

在 Lab4 的代码中, 我们在 `task_init` 的时候创建了用户栈, `load_program` 的时候拷贝了 load segment, 并通过 `create_mapping` 在页表中创建了映射. 在本次实验中, 我们将修改为 demand paging 的方式, 也就是在初始化 task 的时候不进行任何的映射

(除了内核栈以及页表以外也不需要开辟其他空间), 而是在发生缺页异常的时候检测到是记录在 vma 中的合法地址后, 再分配页面并进行映射.

2.4 RISC-V Page Faults

在 RISC-V 中, 当系统运行发生异常时, 可通过解析 scause 寄存器的值, 识别如下三种不同的 page fault:

Interrupt	Exception Code	Description
0	12	Instruction Page Fault
0	13	Load Page Fault
0	15	Store/AMO Page Fault

2.5 处理 page fault 的方式

处理缺页异常时可能所需的信息如下:

- 触发 page fault 时访问的虚拟内存地址. 当触发 page fault 时, stval 寄存器被硬件自动设置为该出错的 VA 地址
- 导致 page fault 的类型, 保存在 scause 寄存器中

Interrupt	Exception Code	Description
0	12	Instruction Page Fault
0	13	Load Page Fault
0	15	Store/AMO Page Fault

- 发生 page fault 时的指令执行位置, 保存在 sepc 中
- 当前进程合法的 VMA 映射关系, 保存在 vm_area_struct 链表中
- 发生异常的虚拟地址对应的 PTE (page table entry) 中记录的信息

总的说来, 处理缺页异常需要进行以下步骤:

1. 捕获异常
2. 寻找当前 task 中导致产生了异常的地址对应的 VMA
 - 如果当前访问的虚拟地址在 VMA 中没有记录, 即是不合法的地址, 则运行出错 (本实验不涉及)

- 如果当前访问的虚拟地址在 VMA 中存在记录, 则需要判断产生异常的原因:
 - 如果是匿名区域, 那么开辟一页内存, 然后把这一页映射到产生异常的 task 的页表中
 - 如果不是, 则访问的页是存在数据的 (如代码), 需要从相应位置读取内容, 然后映射到页表中

3. 返回到产生了该缺页异常的那条指令, 并继续执行程序

2.6 Fork 系统调用

Fork 是 Linux 中的重要系统调用, 它的作用是将进行了该系统调用的 task 完整地复制一份, 并加入 Ready Queue. 这样在下一次调度发生时, 调度器就能够发现多了一个 task. 从这时候开始, 新的 task 就可能被正式从 Ready 调度到 Running, 而开始执行了. 需留意, fork 具有以下特点:

- Fork 通过复制当前进程创建一个新的进程, 新进程称为子进程, 而原进程称为父进程
- 子进程和父进程在不同的内存空间上运行
- Fork 成功时, 父进程返回子进程的 PID, 子进程返回 0; 失败时, 父进程返回 -1
- 创建的子 task 需要深拷贝 task_struct, 调整自己的页表, 栈和 CSR 寄存器等信息, 复制一份在用户态会用到的内存信息 (用户态的栈, 程序的代码和数据等), 并且将自己伪装成是一个因为调度而加入了 Ready Queue 的普通程序来等待调度. 在调度发生时, 这个新 task 就像是原本就在等待调度一样, 被调度器选择并调度.

3 实验具体过程与代码实现

3.1 缺页异常处理

3.1.1 实现虚拟内存管理功能

我们实现了下述两个函数:

- find_vma 函数: 实现对 vm_area_struct 的查找

```

1  /*
2  * @mm : current thread's mm_struct
3  * @addr : the va to look up
4  *
5  * @return : the VMA if found or NULL if not found
6  */
7  struct vm_area_struct *find_vma(struct mm_struct *mm,
    uint64_t addr);

```

- do_mmap 函数: 实现 vm_area_struct 的添加

```

1  /*
2  * @mm : current thread's mm_struct
3  * @addr : the suggested va to map
4  * @len : memory size to map
5  * @vm_pgoff : phdr->p_offset
6  * @vm_filesz: phdr->p_filesz
7  * @flags : flags for the new VMA
8  *
9  * @return : start va
10 */
11 uint64_t do_mmap(struct mm_struct *mm, uint64_t addr,
    uint64_t len, uint64_t vm_pgoff, uint64_t vm_filesz,
    uint64_t flags);

```

具体实现如下:

```

1  struct vm_area_struct *find_vma(struct mm_struct *mm, uint64_t
    addr)
2  {
3      struct vm_area_struct *vma = mm->mmap;
4      while (vma != NULL)
5      {
6          if (vma->vm_start <= addr && vma->vm_end > addr)
7          {
8              return vma;

```

```

9         }
10        vma = vma->vm_next;
11    }
12    return vma;
13 }
14 uint64_t do_mmap(struct mm_struct *mm, uint64_t addr, uint64_t
15                len, uint64_t vm_pgoff, uint64_t vm_filesz, uint64_t flags)
16 {
17     struct vm_area_struct *new_vma = (struct vm_area_struct
18                                       *)kalloc();
19     new_vma->vm_mm = mm;
20     new_vma->vm_start = addr;
21     new_vma->vm_end = addr + len;
22     new_vma->vm_flags = flags;
23     new_vma->vm_pgoff = vm_pgoff;
24     new_vma->vm_filesz = vm_filesz;
25     new_vma->vm_next = mm->mmap;
26     new_vma->vm_prev = NULL;
27     if (mm->mmap != NULL)
28     {
29         mm->mmap->vm_prev = new_vma;
30     }
31     mm->mmap = new_vma;
32     Log("construct vma: [%p, %p) with flags %p for mm %p",
33         new_vma->vm_start, new_vma->vm_end, new_vma->vm_flags, mm);
34     return addr;
35 }

```

3.1.2 修改 tas_init

具体实现如下:

```

1 void task_init()
2 {
3     ...
4     for (int i = 0; i < NR_TASKS; i++)

```



```

5      {
6          if (task[i] == NULL)
7          {
8              ...
9
10             do_mmap(&(task[i]->mm), USER_END - PGSIZE, PGSIZE, 0,
11                     0, VM_READ | VM_WRITE | VM_ANON);
12             // load_binary(task[i]->pagetable, _sramdisk,
13                          _eramdisk);
14             task[i]->thread.sepc =
15                 load_elf_lazy(task[i]->pagetable,
16                              _sramdisk, task[i]);
17             break;
18         }
19     }
20     printk("...task_init done!\n");
21 }

```

我们在 `task_init` 的时候创建了用户栈, `load_program` 的时候拷贝了 `load segment`, 并通过 `create_mapping` 在页表中创建了映射. 在本次实验中, 我们将修改为 `demand paging` 的方式, 也就是在初始化 `task` 的时候不进行任何的映射, 仅仅将用户栈和所有所需的用户空间的 `VMA` 添加到 `mm` 中, 不进行内存分配和页表映射.

其中的 `load_elf_lazy` 函数的实现如下:

```

1  uint64_t load_elf_lazy(pagetable_ptr_t pagetable, char *start,
2                        struct task_struct *new_proc)
3  {
4      Elf64_Ehdr *elf = (Elf64_Ehdr *)start;
5      for (int i = 0; i < elf->e_phnum; i++)
6      {
7          Elf64_Phdr *phdr = (Elf64_Phdr *)((uint64_t)start +
8              elf->e_phoff + i * sizeof(Elf64_Phdr));
9          if (phdr->p_type != PT_LOAD)
10             continue;
11         uint64_t flags = 0;
12         if (phdr->p_flags & PF_X)

```

```

11         flags |= PTE_X;
12         if (phdr->p_flags & PF_W)
13             flags |= PTE_W;
14         if (phdr->p_flags & PF_R)
15             flags |= PTE_R;
16         do_mmap(&(new_proc->mm), phdr->p_vaddr, phdr->p_memsz,
17                phdr->p_offset, phdr->p_filesz, flags);
18     }
19     return elf->e_entry;
20 }

```

可以看出, 我们在 `load_elf_lazy` 函数中, 仅仅将用户空间的 VMA 添加到 `mm` 中, 不进行内存分配和页表映射.

3.1.3 实现 page fault handler

按照上文2.5的步骤, 我们实现了 `page fault handler` 函数:

```

1  void trap_handler(uint64_t scause, uint64_t sepc, struct pt_regs
2      *regs)
3  {
4      uint64_t bad_addr;
5
6      if ((scause >> 63) == 0)
7      {
8          ...
9          case 0xC:
10             bad_addr = csr_read(stval);
11             Log("Instruction page fault at %p", bad_addr);
12             do_page_fault(bad_addr, VM_EXEC);
13             break;
14             case 0xD:
15                 bad_addr = csr_read(stval);
16                 Log("Load page fault at %p", bad_addr);
17                 do_page_fault(bad_addr, VM_READ);
18                 break;
19             case 0xF:

```

```

19         bad_addr = csr_read(stval);
20         Log("Store/AMO page fault at %p", bad_addr);
21         do_page_fault(bad_addr, VM_WRITE);
22         break;
23     ...
24 }
25 ...
26 }
27 void do_page_fault(uint64_t bad_addr, uint64_t flags)
28 {
29     struct vm_area_struct *vma = find_vma(&(current->mm),
30         bad_addr);
31     if (vma == NULL)
32     {
33         Err("Page fault at %p out of vma", bad_addr);
34         return;
35     }
36     if ((vma->vm_flags & flags) == 0)
37     {
38         if (flags == VM_EXEC)
39         {
40             Err("Page fault at %p, no exec permission", bad_addr);
41         }
42         else if (flags == VM_READ)
43         {
44             Err("Page fault at %p, no read permission", bad_addr);
45         }
46         else if (flags == VM_WRITE)
47         {
48             Err("Page fault at %p, no write permission", bad_addr);
49         }
50         return;
51     }
52     uint64_t pte_flags = PTE_V | PTE_U;
53     if (vma->vm_flags & VM_EXEC)

```

```

53         pte_flags |= PTE_X;
54     if (vma->vm_flags & VM_WRITE)
55     {
56         ... // COW
57         pte_flags |= PTE_W;
58     }
59     if (vma->vm_flags & VM_READ)
60         pte_flags |= PTE_R;
61     uint64_t new_pg = (uint64_t)kalloc();
62     memset((void *)new_pg, 0, PGSIZE);
63
64     if (!(vma->vm_flags & VM_ANON) && vma->vm_start <= bad_addr
        && vma->vm_end > bad_addr)
65     {
66         uint64_t copy_start = PGROUNDDOWN(bad_addr) <
            vma->vm_start ? vma->vm_start : PGROUNDDOWN(bad_addr);
67         uint64_t copy_end = (PGROUNDDOWN(bad_addr) + PGSIZE) >
            vma->vm_start + vma->vm_filesz ? vma->vm_start +
            vma->vm_filesz : (PGROUNDDOWN(bad_addr) + PGSIZE);
68         if (copy_end > copy_start)
69         {
70             uint64_t copy_size = copy_end - copy_start;
71             uint64_t copy_offset = copy_start - vma->vm_start +
                vma->vm_pgoff;
72             memcpy((void *)new_pg + copy_start -
                PGROUNDDOWN(bad_addr), (void *)(_sramdisk +
                copy_offset), copy_size);
73         }
74     }
75     create_mapping(current->pagetable, PGROUNDDOWN(bad_addr),
        VA2PA(new_pg), PGSIZE, pte_flags);
76     Log("Page fault at %p, create mapping to %p", bad_addr,
        new_pg);
77 }

```

注意中间有一段逻辑是用于处理 COW 暂且不提.

3.2 实现 fork 系统调用

3.2.1 拷贝内核栈

```
1  uint64_t fork(struct pt_regs *old_regs)
2  {
3      int pid;
4      for (pid = 1; pid < NR_TASKS; pid++)
5      {
6          if (task[pid] == NULL)
7          {
8              break;
9          }
10     }
11     if (pid == NR_TASKS)
12     {
13         Err("No more process can be created\n");
14         return -1;
15     }
16     uint64_t kernel_stack = (uint64_t)kalloc();
17     pagetable_ptr_t pagetable = (pagetable_ptr_t)kalloc();
18     memcpy((void *)kernel_stack, (void *)current, PGSIZE);
19     task[pid] = (struct task_struct *)kernel_stack;
20     task[pid]->pid = pid;
21     task[pid]->pagetable = pagetable;
22     task[pid]->mm.mmap = NULL;
23     ...
24 }
```

我们在 fork 函数中首先找到一个空的进程位, 随后我们拷贝内核栈, 分配内核页表, 设置内核序号, 初始化进程 VMA.

3.2.2 创建子进程页表

流程如下:

- 拷贝内核页表 swapper_pg_dir
- 遍历父进程 vma, 并遍历父进程页表
 - 将这个 vma 也添加到新进程的 vma 链表中
 - 如果该 vma 项有对应的页表项存在 (说明已经创建了映射), 则需要深拷贝一整页的内容并映射到新页表中
- 复制父进程的 VMA

具体实现如下:

```

1  void fork(struct pt_regs *old_regs)
2  {
3      ...
4      copy_pgtbl(task[pid]->pagetable,
5                  (pagetable_ptr_t)swapper_pg_dir);
6      vmas_deep_copy(&(task[pid]->mm), &(current->mm),
7                     task[pid]->pagetable, current->pagetable);
8      ...
9  }
10 void vmas_deep_copy(struct mm_struct *new_mm, struct mm_struct
11                     *old_mm, pagetable_ptr_t new_pagetable, pagetable_ptr_t
12                     old_pagetable)
13 {
14     for (struct vm_area_struct *vma = old_mm->mmap; vma != NULL;
15          vma = vma->vm_next)
16     {
17         do_mmap(new_mm, vma->vm_start, vma->vm_end -
18                 vma->vm_start, vma->vm_pgoff, vma->vm_filesz,
19                 vma->vm_flags);
20         vma_pages_copy(new_pagetable, old_pagetable,
21                         PGROUNDDOWN(vma->vm_start), PGROUNDUP(vma->vm_end));
22     }
23 }
24 void vma_pages_copy(pagetable_ptr_t new_pagetable,
25                     pagetable_ptr_t old_pagetable, uint64_t start, uint64_t end)
26 {

```

```

18     for (uint64_t va = start; va < end; va += PGSIZE)
19     {
20         int level = 3;
21         pte_t *pte = walk(old_pagetable, va, 0, &level);
22         if (pte == 0 || !(*pte & PTE_V))
23         {
24             continue;
25         }
26         uint64_t old_pa = PTE2ADDR(*pte, 3);
27         uint64_t old_va = PA2VA(old_pa);
28         uint64_t new_pa = (uint64_t)kalloc();
29         memcpy((void *)new_pa, (void *)old_va, PGSIZE);
30         create_mapping(new_pagetable, va, VA2PA(new_pa), PGSIZE,
31                        PTE2FLAG(*pte));
32     }

```

我们在这里实现了 `vmass_deep_copy` 函数, 用于复制父进程的 VMA, 并在其中调用了 `vma_pages_copy` 函数, 用于复制父进程已分配的所有页到新的进程中, 并映射到新页表中.

3.2.3 处理进程返回逻辑

代码实现如下:

```

1
2 uint64_t fork(struct pt_regs *old_regs)
3 {
4     ...
5     task[pid]->thread.ra = (uint64_t)__ret_from_fork;
6     task[pid]->thread.sp = kernel_stack + PGSIZE - sizeof(struct
7         pt_regs);
8     task[pid]->thread.sscratch = 0;
9     struct pt_regs *regs = (struct pt_regs *) (kernel_stack + PGSIZE
10        - sizeof(struct pt_regs));
11     regs->x[10] = 0;
12     regs->sepc += 4;

```

```

11     Log("fork: pid = %d", pid);
12     return pid;
13 }

```

我们这里直接设置新进程的返回地址为 `call trap_handler` 的返回地址, 因为在这里我们的保存和恢复函数只使用了 `sp`, 并且退出时的 `sp` 十分好计算即内核栈顶减去一个 `pt_regs` 的大小, 这样我们在配置子进程的内核上下文时, 只需维护 `sp` 即可. 同时还需要设置新进程的 `sscratch` 为 0, 因为处于内核态. 此外我们还需要设置 `x[10]` 为 0, 这是因为子进程 `fork` 的返回值为 0. 最后我们将 `sepc` 加 4, 这是系统调用的正常操作, 不然会重复执行 `fork`.

3.2.4 写时复制 COW

接下来在我们 `do_fork` 创建页面, 拷贝内容, 创建页表的时候, 只需要:

- 将物理页的引用计数加一
- 将父进程的该地址对应的页表项的 `PTE_W` 位置 0
- 注意因为修改了页表项权限, 所以全部修改完成后需要通过 `sfence.vma` 刷新 TLB
- 为子进程创建一个新的页表项, 指向父进程的物理页, 且权限不带 `PTE_W`

这样在父子进程想要写入的时候, 就会触发 `page fault`, 然后再由我们在 `page fault handler` 中进行 COW. 在 `handler` 中, 我们只需要判断, 如果发生了写错误, 且 `vma` 的 `VM_WRITE` 位为 1, 而且对应地址有 `pte`(进行了映射) 但 `pte` 的 `PTE_W` 位为 0, 那么就可以断定这是一个写时复制的页面, 我们只需要在这个时候拷贝一份原来的页面, 重新创建一个映射即可.

我们首先实现了 `vmaz_lazy_copy` 函数, 用于在 `fork` 时进行拷贝:

```

1     void vmaz_lazy_copy(struct mm_struct *new_mm, struct mm_struct
      *old_mm, pagetable_ptr_t new_pagetable, pagetable_ptr_t
      old_pagetable)
2     {
3         for (struct vm_area_struct *vma = old_mm->mmap; vma != NULL;
      vma = vma->vm_next)
4         {
5             do_mmap(new_mm, vma->vm_start, vma->vm_end -
      vma->vm_start, vma->vm_pgoff, vma->vm_filesz,
      vma->vm_flags);

```



```

6         vma_pages_lazy_copy(new_pagetable, old_pagetable,
          PGROUNDDOWN(vma->vm_start), PGROUNDUP(vma->vm_end));
7     }
8     flush_tlb();
9 }
10 void vma_pages_lazy_copy(pagetable_ptr_t new_pagetable,
    pagetable_ptr_t old_pagetable, uint64_t start, uint64_t end)
11 {
12     for (uint64_t va = start; va < end; va += PGSIZE)
13     {
14         int level = 3;
15         pte_t *pte = walk(old_pagetable, va, 0, &level);
16         if (pte == 0 || !(*pte & PTE_V))
17         {
18             continue;
19         }
20         if ((*pte & PTE_W))
21         {
22             *pte &= ~PTE_W;
23         }
24         get_page((void *)PA2VA(PTE2ADDR(*pte, 3)));
25         create_mapping(new_pagetable, va, PTE2ADDR(*pte, 3),
            PGSIZE, PTE2FLAG(*pte));
26     }
27 }

```

在这个函数中, 我们首先遍历父进程的 VMA, 并在新进程中创建相同的 VMA, 然后遍历父进程的页表, 并在新进程的页表中创建相同的映射, 并将物理页的引用计数加一, 若发现 PTE_W 位为 1, 则将两张页表项的 PTE_W 位清零. 注意在修改页表项权限后, 我们需要通过 flush_tlb 刷新 TLB.

我们接着实现了 COW trap 处理逻辑, 我们首先捕获到 COW, 然后通过 handle_COW 函数处理:

```

1     void do_page_fault(uint64_t bad_addr, uint64_t flags)
2     {
3         ...

```

```

4      if (vma->vm_flags & VM_WRITE)
5      {
6          if (flags == VM_WRITE)
7          {
8              int level = 3;
9              pte_t *pte = walk(current->pagetable, bad_addr, 0,
10                             &level);
11              if (pte && !(*pte & PTE_W) && (*pte & PTE_V))
12              {
13                  handle_COW(pte);
14                  Log("handle COW for process %d at address %p",
15                      current->pid, bad_addr);
16                  return;
17              }
18          }
19          pte_flags |= PTE_W;
20      }
21      ...
22  }
23  void handle_COW(pte_t *pte)
24  {
25      uint64_t old_va = PA2VA(PTE2ADDR(*pte, 3));
26      if (get_page_refcnt((void *)old_va) > 1)
27      {
28          uint64_t new_va = (uint64_t)kalloc();
29          memcpy((void *)new_va, (void *)old_va, PGSIZE);
30          put_page((void *)old_va);
31          *pte = ADDR2PTE(VA2PA(new_va), PTE2FLAG(*pte), 3);
32      }
33      *pte |= PTE_W;
34      flush_tlb();
35  }

```

捕获 COW 的逻辑是, 如果发生了写错误, 且 vma 的 VM_WRITE 位为 1, 而且对应地址有 pte(进行了映射) 但 pte 的 PTE_W 位为 0, 那么就可以断定这是一个写

时复制的页面.

在 `handle_COW` 函数中, 我们首先判断该物理页的引用计数是否大于 1, 若是则说明有多个进程共享这个物理页, 我们需要为这个物理页创建一个新的物理页, 并将原物理页的内容拷贝到新的物理页中, 否则直接用原物理页即可, 然后将原物理页的引用计数减一, 并将新的物理页映射到页表中, 并将 `PTE_W` 位置 1, 最后刷新 TLB.

4 实验结果与分析

4.1 测试缺页处理

我们首先运行 PFH1 测试:

```

...buddy_init done!
...mm_init done!
[vm.c,122,create_mapping] create_mapping: [0xffffffff00020000, 0xffffffff00020500] -> [0x8020000, 0x8020500] with perm 0xb
[vm.c,122,create_mapping] create_mapping: [0xffffffff00020500, 0xffffffff00020600] -> [0x8020500, 0x8020600] with perm 0x3
[vm.c,122,create_mapping] create_mapping: [0xffffffff00020600, 0xffffffff00020700] -> [0x8020600, 0x8020700] with perm 0x7
[vm.c,287,do_mmap] construct vma: [0x3ffffff00, 0x400000000] with flags 0x7 for mm 0xffffffff00030e0b0
[vm.c,287,do_mmap] construct vma: [0x100e8, 0x12620] with flags 0xe for mm 0xffffffff00030e0b0
[vm.c,287,do_mmap] construct vma: [0x3ffffff00, 0x400000000] with flags 0x7 for mm 0xffffffff0003520b0
[vm.c,287,do_mmap] construct vma: [0x100e8, 0x12620] with flags 0xe for mm 0xffffffff0003520b0
[vm.c,287,do_mmap] construct vma: [0x3ffffff00, 0x400000000] with flags 0x7 for mm 0xffffffff0003960b0
[vm.c,287,do_mmap] construct vma: [0x100e8, 0x12620] with flags 0xe for mm 0xffffffff0003960b0
[vm.c,287,do_mmap] construct vma: [0x3ffffff00, 0x400000000] with flags 0x7 for mm 0xffffffff0003da0b0
[vm.c,287,do_mmap] construct vma: [0x100e8, 0x12620] with flags 0xe for mm 0xffffffff0003da0b0
...task_init done!
switch to [PID = 2 PRIORITY = 10 COUNTER = 10]
[trap.c,112,trap_handler] Instruction page fault at 0x100e8
[vm.c,122,create_mapping] create_mapping: [0x10000, 0x11000] -> [0x8041e000, 0x8041f000] with perm 0x1f
[trap.c,237,do_page_fault] Page fault at 0x100e8, create mapping to 0xffffffff00041e000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x100e8
[trap.c,124,trap_handler] Store/AMO page fault at 0x3ffffff00
[vm.c,122,create_mapping] create_mapping: [0x3ffffff00, 0x400000000] -> [0x80421000, 0x80422000] with perm 0x17
[trap.c,237,do_page_fault] Page fault at 0x3ffffff00, create mapping to 0xffffffff000421000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x10178
[trap.c,118,trap_handler] Load page fault at 0x12230
[vm.c,122,create_mapping] create_mapping: [0x12000, 0x13000] -> [0x80424000, 0x80425000] with perm 0x1f
[trap.c,237,do_page_fault] Page fault at 0x12230, create mapping to 0xffffffff000424000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x10198
[trap.c,112,trap_handler] Instruction page fault at 0x110ac
[vm.c,122,create_mapping] create_mapping: [0x11000, 0x12000] -> [0x80425000, 0x80426000] with perm 0x1f
[trap.c,237,do_page_fault] Page fault at 0x110ac, create mapping to 0xffffffff000425000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x110ac
[U-MODE] pid: 2, sp is 0x3ffffffe0, this is print No.1
switch to [PID = 1 PRIORITY = 7 COUNTER = 7]
[trap.c,112,trap_handler] Instruction page fault at 0x100e8
[vm.c,122,create_mapping] create_mapping: [0x10000, 0x11000] -> [0x80426000, 0x80427000] with perm 0x1f
[trap.c,237,do_page_fault] Page fault at 0x100e8, create mapping to 0xffffffff000426000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x100e8
[trap.c,124,trap_handler] Store/AMO page fault at 0x3ffffff00
[vm.c,122,create_mapping] create_mapping: [0x3ffffff00, 0x400000000] -> [0x80429000, 0x8042a000] with perm 0x17
[trap.c,237,do_page_fault] Page fault at 0x3ffffff00, create mapping to 0xffffffff000429000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x10178
[trap.c,118,trap_handler] Load page fault at 0x12230
[vm.c,122,create_mapping] create_mapping: [0x12000, 0x13000] -> [0x8042c000, 0x8042d000] with perm 0x1f
[trap.c,237,do_page_fault] Page fault at 0x12230, create mapping to 0xffffffff00042c000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x10198
[trap.c,112,trap_handler] Instruction page fault at 0x110ac
[vm.c,122,create_mapping] create_mapping: [0x11000, 0x12000] -> [0x8042d000, 0x8042e000] with perm 0x1f
[trap.c,237,do_page_fault] Page fault at 0x110ac, create mapping to 0xffffffff00042d000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x110ac
[U-MODE] pid: 1, sp is 0x3ffffffe0, this is print No.1
switch to [PID = 3 PRIORITY = 4 COUNTER = 4]
[trap.c,112,trap_handler] Instruction page fault at 0x100e8
[vm.c,122,create_mapping] create_mapping: [0x10000, 0x11000] -> [0x8042e000, 0x8042f000] with perm 0x1f

```

图 1: make run TEST=PFH1 image-1

```

[trap.c,163,trap_handler] Trap occurs at sepc = 0x110ac
[U-MODE] pid: 1, sp is 0x3fffffff0, this is print No.1
switch to [PID = 3 PRIORITY = 4 COUNTER = 4]
[trap.c,112,trap_handler] Instruction page fault at 0x100e8
[vm.c,122,create_mapping] create_mapping: [0x10000, 0x11000) -> [0x8042e000, 0x8042f000) with perm 0x1f
[trap.c,237,do_page_fault] Page fault at 0x100e8, create mapping to 0xffffffff00042e000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x100e8
[trap.c,124,trap_handler] Store/AMO page fault at 0x3fffffff8
[vm.c,122,create_mapping] create_mapping: [0x3fffffff000, 0x4000000000) -> [0x80431000, 0x80432000) with perm
0x17
[trap.c,237,do_page_fault] Page fault at 0x3fffffff8, create mapping to 0xffffffff000431000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x10178
[trap.c,118,trap_handler] Load page fault at 0x12230
[vm.c,122,create_mapping] create_mapping: [0x12000, 0x13000) -> [0x80434000, 0x80435000) with perm 0x1f
[trap.c,237,do_page_fault] Page fault at 0x12230, create mapping to 0xffffffff000434000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x10198
[trap.c,112,trap_handler] Instruction page fault at 0x110ac
[vm.c,122,create_mapping] create_mapping: [0x11000, 0x12000) -> [0x80435000, 0x80436000) with perm 0x1f
[trap.c,237,do_page_fault] Page fault at 0x110ac, create mapping to 0xffffffff000435000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x110ac
[U-MODE] pid: 3, sp is 0x3fffffff0, this is print No.1
switch to [PID = 4 PRIORITY = 1 COUNTER = 1]
[trap.c,112,trap_handler] Instruction page fault at 0x100e8
[vm.c,122,create_mapping] create_mapping: [0x10000, 0x11000) -> [0x80436000, 0x80437000) with perm 0x1f
[trap.c,237,do_page_fault] Page fault at 0x100e8, create mapping to 0xffffffff000436000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x100e8
[trap.c,124,trap_handler] Store/AMO page fault at 0x3fffffff8
[vm.c,122,create_mapping] create_mapping: [0x3fffffff000, 0x4000000000) -> [0x80439000, 0x8043a000) with perm
0x17
[trap.c,237,do_page_fault] Page fault at 0x3fffffff8, create mapping to 0xffffffff000439000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x10178
[trap.c,118,trap_handler] Load page fault at 0x12230
[vm.c,122,create_mapping] create_mapping: [0x12000, 0x13000) -> [0x8043c000, 0x8043d000) with perm 0x1f
[trap.c,237,do_page_fault] Page fault at 0x12230, create mapping to 0xffffffff00043c000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x10198
[trap.c,112,trap_handler] Instruction page fault at 0x110ac
[vm.c,122,create_mapping] create_mapping: [0x11000, 0x12000) -> [0x8043d000, 0x8043e000) with perm 0x1f
[trap.c,237,do_page_fault] Page fault at 0x110ac, create mapping to 0xffffffff00043d000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x110ac
[U-MODE] pid: 4, sp is 0x3fffffff0, this is print No.1
switch to [PID = 2 PRIORITY = 10 COUNTER = 10]
[U-MODE] pid: 2, sp is 0x3fffffff0, this is print No.2
switch to [PID = 1 PRIORITY = 7 COUNTER = 7]
switch to [PID = 3 PRIORITY = 4 COUNTER = 4]
switch to [PID = 4 PRIORITY = 1 COUNTER = 1]
switch to [PID = 2 PRIORITY = 10 COUNTER = 10]
switch to [PID = 1 PRIORITY = 7 COUNTER = 7]
[U-MODE] pid: 1, sp is 0x3fffffff0, this is print No.2
switch to [PID = 3 PRIORITY = 4 COUNTER = 4]
switch to [PID = 4 PRIORITY = 1 COUNTER = 1]
switch to [PID = 2 PRIORITY = 10 COUNTER = 10]
[U-MODE] pid: 2, sp is 0x3fffffff0, this is print No.3
switch to [PID = 1 PRIORITY = 7 COUNTER = 7]
switch to [PID = 3 PRIORITY = 4 COUNTER = 4]
switch to [PID = 4 PRIORITY = 1 COUNTER = 1]

```

图 2: make run TEST=PFH1 image-2

可以看到所有的页错误都被正确处理了, 在运行中分配了新的页. 并且之后正常执行了.

我们再运行 PFH2 测试:

```

...buddy_init done!
...mm_init done!
[vm.c,122,create_mapping] create_mapping: [0xffffffff00020000, 0xffffffff00020500] -> [0x8020000, 0x8020500] with perm 0xb
[vm.c,122,create_mapping] create_mapping: [0xffffffff00020500, 0xffffffff00020600] -> [0x8020500, 0x8020600] with perm 0x3
[vm.c,122,create_mapping] create_mapping: [0xffffffff00020600, 0xffffffff00000000] -> [0x8020600, 0x8000000] with perm 0x7
[vm.c,287,do_mmap] construct vma: [0x3ffffff00, 0x400000000] with flags 0x7 for mm 0xffffffff00030e0b0
[vm.c,287,do_mmap] construct vma: [0x100e8, 0x135e8] with flags 0xe for mm 0xffffffff00030e0b0
[vm.c,287,do_mmap] construct vma: [0x3ffffff00, 0x400000000] with flags 0x7 for mm 0xffffffff0003520b0
[vm.c,287,do_mmap] construct vma: [0x100e8, 0x135e8] with flags 0xe for mm 0xffffffff0003520b0
[vm.c,287,do_mmap] construct vma: [0x3ffffff00, 0x400000000] with flags 0x7 for mm 0xffffffff0003960b0
[vm.c,287,do_mmap] construct vma: [0x100e8, 0x135e8] with flags 0xe for mm 0xffffffff0003960b0
[vm.c,287,do_mmap] construct vma: [0x3ffffff00, 0x400000000] with flags 0x7 for mm 0xffffffff0003da0b0
[vm.c,287,do_mmap] construct vma: [0x100e8, 0x135e8] with flags 0xe for mm 0xffffffff0003da0b0
...task_init done!
switch to [PID = 2 PRIORITY = 10 COUNTER = 10]
[trap.c,112,trap_handler] Instruction page fault at 0x100e8
[vm.c,122,create_mapping] create_mapping: [0x10000, 0x11000] -> [0x8041e000, 0x8041f000] with perm 0x1f
[trap.c,237,do_page_fault] Page fault at 0x100e8, create mapping to 0xffffffff00041e000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x100e8
[trap.c,124,trap_handler] Store/AMO page fault at 0x3ffffff8
[vm.c,122,create_mapping] create_mapping: [0x3ffffff00, 0x400000000] -> [0x80421000, 0x80422000] with perm 0x17
[trap.c,237,do_page_fault] Page fault at 0x3ffffff8, create mapping to 0xffffffff000421000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x10178
[trap.c,118,trap_handler] Load page fault at 0x131f0
[vm.c,122,create_mapping] create_mapping: [0x13000, 0x14000] -> [0x80424000, 0x80425000] with perm 0x1f
[trap.c,237,do_page_fault] Page fault at 0x131f0, create mapping to 0xffffffff000424000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x10194
[trap.c,112,trap_handler] Instruction page fault at 0x1107c
[vm.c,122,create_mapping] create_mapping: [0x11000, 0x12000] -> [0x80425000, 0x80426000] with perm 0x1f
[trap.c,237,do_page_fault] Page fault at 0x1107c, create mapping to 0xffffffff000425000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x1107c
[U-MODE] pid: 2, increment: 0
switch to [PID = 1 PRIORITY = 7 COUNTER = 7]
[trap.c,112,trap_handler] Instruction page fault at 0x100e8
[vm.c,122,create_mapping] create_mapping: [0x10000, 0x11000] -> [0x80426000, 0x80427000] with perm 0x1f
[trap.c,237,do_page_fault] Page fault at 0x100e8, create mapping to 0xffffffff000426000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x100e8
[trap.c,124,trap_handler] Store/AMO page fault at 0x3ffffff8
[vm.c,122,create_mapping] create_mapping: [0x3ffffff00, 0x400000000] -> [0x80429000, 0x8042a000] with perm 0x17
[trap.c,237,do_page_fault] Page fault at 0x3ffffff8, create mapping to 0xffffffff000429000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x10178
[trap.c,118,trap_handler] Load page fault at 0x131f0
[vm.c,122,create_mapping] create_mapping: [0x13000, 0x14000] -> [0x8042c000, 0x8042d000] with perm 0x1f
[trap.c,237,do_page_fault] Page fault at 0x131f0, create mapping to 0xffffffff00042c000

```

图 3: make run TEST=PFH2

可以看到页 [0x12000, 0x13000] 没有被映射。

两个测试均符合预期。

4.2 测试 fork

我们运行 FORK1 测试:

```

...buddy_init done!
...mm_init done!
[vm.c,122,create_mapping] create_mapping: [0xffffffff00020000, 0xffffffff000205000] -> [0x80200000, 0x80205000] with perm 0xb
[vm.c,122,create_mapping] create_mapping: [0xffffffff000205000, 0xffffffff000206000] -> [0x80205000, 0x80206000] with perm 0x3
[vm.c,122,create_mapping] create_mapping: [0xffffffff000206000, 0xffffffff000800000] -> [0x80206000, 0x80800000] with perm 0x7
[vm.c,287,do_mmap] construct vma: [0x3fffffff000, 0x4000000000] with flags 0x7 for mm 0xffffffff00030e0b0
[vm.c,287,do_mmap] construct vma: [0x100e8, 0x126c0] with flags 0xe for mm 0xffffffff00030e0b0
...task_init done!
switch to [PID = 1 PRIORITY = 7 COUNTER = 7]
[trap.c,112,trap_handler] Instruction page fault at 0x100e8
[vm.c,122,create_mapping] create_mapping: [0x10000, 0x11000] -> [0x80352000, 0x80353000] with perm 0x1f
[trap.c,237,do_page_fault] Page fault at 0x100e8, create mapping to 0xffffffff000352000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x100e8
[trap.c,124,trap_handler] Store/AMO page fault at 0x3fffffff8
[vm.c,122,create_mapping] create_mapping: [0x3fffffff000, 0x4000000000] -> [0x80355000, 0x80356000] with perm 0x17
[trap.c,237,do_page_fault] Page fault at 0x3fffffff8, create mapping to 0xffffffff000355000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x101ac
[vm.c,287,do_mmap] construct vma: [0x100e8, 0x126c0] with flags 0xe for mm 0xffffffff0003580b0
[vm.c,122,create_mapping] create_mapping: [0x10000, 0x11000] -> [0x80352000, 0x80353000] with perm 0x5b
[vm.c,287,do_mmap] construct vma: [0x3fffffff000, 0x4000000000] with flags 0x7 for mm 0xffffffff0003580b0
[vm.c,122,create_mapping] create_mapping: [0x3fffffff000, 0x4000000000] -> [0x80355000, 0x80356000] with perm 0xd3
[proc.c,234,fork] fork: pid = 2
[trap.c,124,trap_handler] Store/AMO page fault at 0x3fffffff8
[trap.c,215,do_page_fault] handle COW for process 1 at address 0x3fffffff8
[trap.c,163,trap_handler] Trap occurs at sepc = 0x1013c
[trap.c,118,trap_handler] Load page fault at 0x122d0
[vm.c,122,create_mapping] create_mapping: [0x12000, 0x13000] -> [0x803a1000, 0x803a2000] with perm 0x1f
[trap.c,237,do_page_fault] Page fault at 0x122d0, create mapping to 0xffffffff0003a1000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x10228
[trap.c,112,trap_handler] Instruction page fault at 0x11114
[vm.c,122,create_mapping] create_mapping: [0x11000, 0x12000] -> [0x803a2000, 0x803a3000] with perm 0x1f
[trap.c,237,do_page_fault] Page fault at 0x11114, create mapping to 0xffffffff0003a2000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x11114
[U-PARENT] pid: 1 is running! global_variable: 0
switch to [PID = 2 PRIORITY = 7 COUNTER = 7]
[trap.c,124,trap_handler] Store/AMO page fault at 0x3fffffff8
[trap.c,215,do_page_fault] handle COW for process 2 at address 0x3fffffff8
[trap.c,163,trap_handler] Trap occurs at sepc = 0x1013c
[trap.c,118,trap_handler] Load page fault at 0x122d0
[vm.c,122,create_mapping] create_mapping: [0x12000, 0x13000] -> [0x803a3000, 0x803a4000] with perm 0x1f
[trap.c,237,do_page_fault] Page fault at 0x122d0, create mapping to 0xffffffff0003a3000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x101e0
[trap.c,112,trap_handler] Instruction page fault at 0x11114
[vm.c,122,create_mapping] create_mapping: [0x11000, 0x12000] -> [0x803a4000, 0x803a5000] with perm 0x1f
[trap.c,237,do_page_fault] Page fault at 0x11114, create mapping to 0xffffffff0003a4000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x11114
[U-CHILD] pid: 2 is running! global_variable: 0
switch to [PID = 1 PRIORITY = 7 COUNTER = 7]
switch to [PID = 2 PRIORITY = 7 COUNTER = 7]
switch to [PID = 1 PRIORITY = 7 COUNTER = 7]
switch to [PID = 2 PRIORITY = 7 COUNTER = 7]
switch to [PID = 1 PRIORITY = 7 COUNTER = 7]
[U-PARENT] pid: 1 is running! global_variable: 1
switch to [PID = 2 PRIORITY = 7 COUNTER = 7]
[U-CHILD] pid: 2 is running! global_variable: 1
switch to [PID = 1 PRIORITY = 7 COUNTER = 7]
switch to [PID = 2 PRIORITY = 7 COUNTER = 7]
switch to [PID = 1 PRIORITY = 7 COUNTER = 7]
switch to [PID = 2 PRIORITY = 7 COUNTER = 7]
switch to [PID = 1 PRIORITY = 7 COUNTER = 7]
switch to [PID = 2 PRIORITY = 7 COUNTER = 7]
switch to [PID = 1 PRIORITY = 7 COUNTER = 7]
[U-PARENT] pid: 1 is running! global_variable: 2
switch to [PID = 2 PRIORITY = 7 COUNTER = 7]
[U-CHILD] pid: 2 is running! global_variable: 2
switch to [PID = 1 PRIORITY = 7 COUNTER = 7]
switch to [PID = 2 PRIORITY = 7 COUNTER = 7]

```

图 4: make run TEST=FORK1

可以看到 global_variable 的值互不影响, 后续 page fault 也是各自为自己的页表添加映射。

我们再运行 FORK2 测试:


```

switch to [PID = 1 PRIORITY = 7 COUNTER = 7]
[trap.c,112,trap_handler] Instruction page fault at 0x100e8
[vm.c,122,create_mapping] create_mapping: [0x10000, 0x11000) -> [0x80352000, 0x80353000) with perm 0x1f
[trap.c,237,do_page_fault] Page fault at 0x100e8, create mapping to 0xffffffff000352000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x100e8
[trap.c,124,trap_handler] Store/AMO page fault at 0x3fffffff8
[vm.c,122,create_mapping] create_mapping: [0x3fffffff000, 0x4000000000) -> [0x80355000, 0x80356000) with perm
0x17
[trap.c,237,do_page_fault] Page fault at 0x3fffffff8, create mapping to 0xffffffff000355000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x101ac
[trap.c,118,trap_handler] Load page fault at 0x12518
[vm.c,122,create_mapping] create_mapping: [0x12000, 0x13000) -> [0x80358000, 0x80359000) with perm 0x1f
[trap.c,237,do_page_fault] Page fault at 0x12518, create mapping to 0xffffffff000358000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x101d0
[trap.c,112,trap_handler] Instruction page fault at 0x112cc
[vm.c,122,create_mapping] create_mapping: [0x11000, 0x12000) -> [0x80359000, 0x8035a000) with perm 0x1f
[trap.c,237,do_page_fault] Page fault at 0x112cc, create mapping to 0xffffffff000359000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x112cc
[trap.c,118,trap_handler] Load page fault at 0x14520
[vm.c,122,create_mapping] create_mapping: [0x14000, 0x15000) -> [0x8035a000, 0x8035b000) with perm 0x1f
[trap.c,237,do_page_fault] Page fault at 0x14520, create mapping to 0xffffffff00035a000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x10434
[U] pid: 1 is running! global_variable: 0
[U] pid: 1 is running! global_variable: 1
[U] pid: 1 is running! global_variable: 2
[trap.c,124,trap_handler] Store/AMO page fault at 0x13520
[vm.c,122,create_mapping] create_mapping: [0x13000, 0x14000) -> [0x8035b000, 0x8035c000) with perm 0x1f
[trap.c,237,do_page_fault] Page fault at 0x13520, create mapping to 0xffffffff00035b000
[trap.c,163,trap_handler] Trap occurs at sepc = 0x10228
[vm.c,287,do_mmap] construct vma: [0x100e8, 0x14910) with flags 0xe for mm 0xffffffff00035c0b0
[vm.c,122,create_mapping] create_mapping: [0x10000, 0x11000) -> [0x80352000, 0x80353000) with perm 0x5b
[vm.c,122,create_mapping] create_mapping: [0x11000, 0x12000) -> [0x80359000, 0x8035a000) with perm 0x5b
[vm.c,122,create_mapping] create_mapping: [0x12000, 0x13000) -> [0x80358000, 0x80359000) with perm 0x5b
[vm.c,122,create_mapping] create_mapping: [0x13000, 0x14000) -> [0x8035b000, 0x8035c000) with perm 0xdb
[vm.c,122,create_mapping] create_mapping: [0x14000, 0x15000) -> [0x8035a000, 0x8035b000) with perm 0xdb
[vm.c,287,do_mmap] construct vma: [0x3fffffff000, 0x4000000000) with flags 0x7 for mm 0xffffffff00035c0b0
[vm.c,122,create_mapping] create_mapping: [0x3fffffff000, 0x4000000000) -> [0x80355000, 0x80356000) with perm
0xd3
[proc.c,234,fork] fork: pid = 2
[trap.c,124,trap_handler] Store/AMO page fault at 0x3fffffff8
[trap.c,215,do_page_fault] handle COW for process 1 at address 0x3fffffff8
[trap.c,163,trap_handler] Trap occurs at sepc = 0x1013c
[trap.c,124,trap_handler] Store/AMO page fault at 0x14520
[trap.c,215,do_page_fault] handle COW for process 1 at address 0x14520
[trap.c,163,trap_handler] Trap occurs at sepc = 0x10448
[U-PARENT] pid: 1 is running! Message: ZJU OS Lab5
[trap.c,124,trap_handler] Store/AMO page fault at 0x12518
[trap.c,215,do_page_fault] handle COW for process 1 at address 0x12518
[trap.c,163,trap_handler] Trap occurs at sepc = 0x103f4
[U-PARENT] pid: 1 is running! global_variable: 3
switch to [PID = 2 PRIORITY = 7 COUNTER = 7]
[trap.c,124,trap_handler] Store/AMO page fault at 0x3fffffff8
[trap.c,215,do_page_fault] handle COW for process 2 at address 0x3fffffff8
[trap.c,163,trap_handler] Trap occurs at sepc = 0x1013c
[trap.c,124,trap_handler] Store/AMO page fault at 0x14520
[trap.c,215,do_page_fault] handle COW for process 2 at address 0x14520
[trap.c,163,trap_handler] Trap occurs at sepc = 0x10448
[U-CHILD] pid: 2 is running! Message: ZJU OS Lab5
[trap.c,124,trap_handler] Store/AMO page fault at 0x12518
[trap.c,215,do_page_fault] handle COW for process 2 at address 0x12518
[trap.c,163,trap_handler] Trap occurs at sepc = 0x1038c
[U-CHILD] pid: 2 is running! global_variable: 3
switch to [PID = 1 PRIORITY = 7 COUNTER = 7]
switch to [PID = 2 PRIORITY = 7 COUNTER = 7]
switch to [PID = 1 PRIORITY = 7 COUNTER = 7]
switch to [PID = 2 PRIORITY = 7 COUNTER = 7]
switch to [PID = 1 PRIORITY = 7 COUNTER = 7]
[U-PARENT] pid: 1 is running! global_variable: 4
switch to [PID = 2 PRIORITY = 7 COUNTER = 7]
[U-CHILD] pid: 2 is running! global_variable: 4
switch to [PID = 1 PRIORITY = 7 COUNTER = 7]
switch to [PID = 2 PRIORITY = 7 COUNTER = 7]

```

图 5: make run TEST=FORK2

本测试的主要输出现象为, 父进程在给 global_variable 自增了三次, 为 placeholder 中赋值了字符串之后才 fork 出子进程, 子进程保留这些信息. PID 2 开始运行时也应该正确输出 ZJU OS Lab5 字符串, 并且 global_variable 从 3 开始自增, 且后续和父进程

互不影响, 符合预期.

最后我们运行 FORK3 测试, 测试结果在下文思考题6.1中.

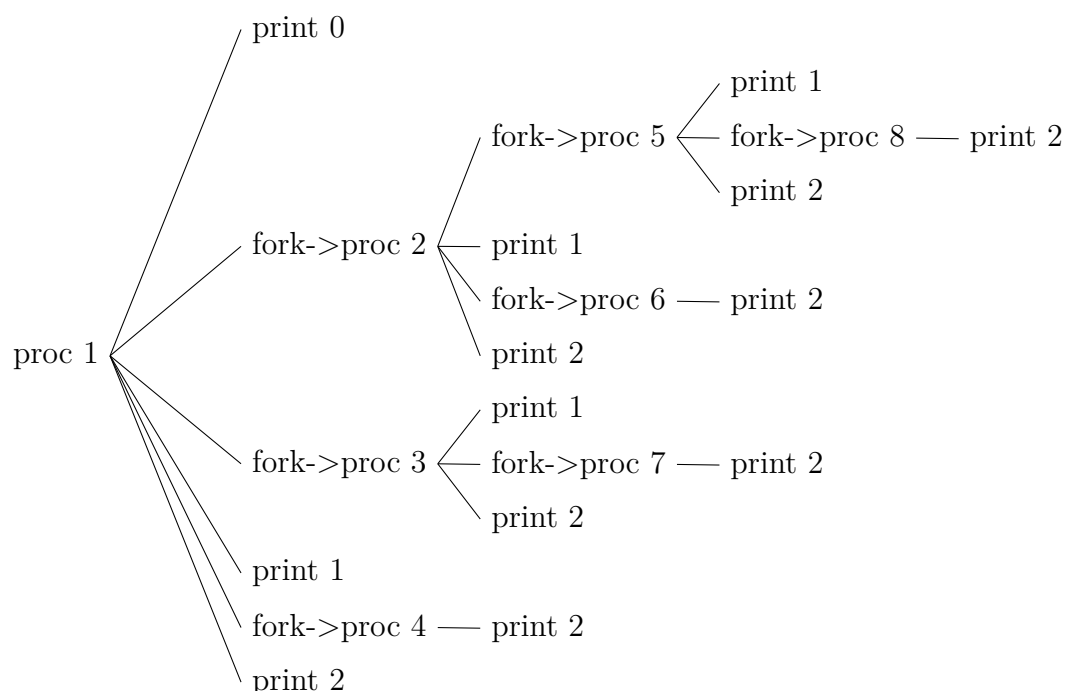
5 遇到的问题及解决方法

问题: 在 FORK3 测试中, 子进程不是从 3 开始输出
解决方法: 一开始我还以为是
因为 COW 没有正确处理双引用页和单引用页, 结果发现是因为没有刷新 TLB.

6 总结与心得

本次实验与上两次实验相比较为简单, 这是因为本次实验的 debug 难度大大降低.
不过我感到有迷惑的一点是, 目前引入了 COW 的页引用机制后, buddy system 只能支持单页分配和释放了, 因为页引用只能支持单页, 混用的话很可能会出错. 那这样的话, 我们引入 buddy system 的目的是什么? 还不如只用链表.

6.1 画图分析 make run TEST=FORK3 的进程 fork 过程, 并呈现出各个进程的 global_variable 应该从几开始输出, 再与你的输出进行对比验证.



由于我们时钟周期设置较大, 我们可以确保在第一个时钟周期内执行完 WAIT 前的指令, 并且我们的调度策略是同优先级序列号越小优先级越高, 所以我们的进程是按照序号从小到大的顺序执行的. 据此, 我们可以画出上述的进程树.

我们可以看出各个进程第一个输出的 global_variable 如下:

proc 1	proc 2	proc 3	proc 4	proc 5	proc 6	proc 7	proc 8
0	1	1	2	1	2	2	2

实际测试如下图:

```

[trap.c,163,trap_handler] Trap occurs at sepc = 0x11130
[U] pid: 1 is running! global_variable: 0
[vm.c,287,do_mmap] construct vma: [0x100e8, 0x126a0) with flags 0xe for mm 0xffffffffe0035a0b0
[vm.c,122,create_mapping] create_mapping: [0x10000, 0x11000) -> [0x80352000, 0x80353000) with perm 0x5b
[vm.c,122,create_mapping] create_mapping: [0x11000, 0x12000) -> [0x80359000, 0x8035a000) with perm 0x5b
[vm.c,122,create_mapping] create_mapping: [0x12000, 0x13000) -> [0x80358000, 0x80359000) with perm 0xdb
[vm.c,287,do_mmap] construct vma: [0x3ffffff00, 0x400000000) with flags 0x7 for mm 0xffffffffe0035a0b0
[vm.c,122,create_mapping] create_mapping: [0x3ffffff00, 0x400000000) -> [0x80355000, 0x80356000) with perm
0xd3
[proc.c,235,fork] fork: pid = 2
[trap.c,124,trap_handler] Store/AMO page fault at 0x3fffffffd8
[trap.c,215,do_page_fault] handle COW for process 1 at address 0x3fffffffd8
[trap.c,163,trap_handler] Trap occurs at sepc = 0x1013c
[vm.c,287,do_mmap] construct vma: [0x100e8, 0x126a0) with flags 0xe for mm 0xffffffffe003a30b0
[vm.c,122,create_mapping] create_mapping: [0x10000, 0x11000) -> [0x80352000, 0x80353000) with perm 0x5b
[vm.c,122,create_mapping] create_mapping: [0x11000, 0x12000) -> [0x80359000, 0x8035a000) with perm 0x5b
[vm.c,122,create_mapping] create_mapping: [0x12000, 0x13000) -> [0x80358000, 0x80359000) with perm 0xdb
[vm.c,287,do_mmap] construct vma: [0x3ffffff00, 0x400000000) with flags 0x7 for mm 0xffffffffe003a30b0
[vm.c,122,create_mapping] create_mapping: [0x3ffffff00, 0x400000000) -> [0x803a2000, 0x803a3000) with perm
0xd3
[proc.c,235,fork] fork: pid = 3
[trap.c,124,trap_handler] Store/AMO page fault at 0x3fffffffd8
[trap.c,215,do_page_fault] handle COW for process 1 at address 0x3fffffffd8
[trap.c,163,trap_handler] Trap occurs at sepc = 0x1013c
[trap.c,124,trap_handler] Store/AMO page fault at 0x122b0
[trap.c,215,do_page_fault] handle COW for process 1 at address 0x122b0
[trap.c,163,trap_handler] Trap occurs at sepc = 0x1021c
[U] pid: 1 is running! global_variable: 1
[vm.c,287,do_mmap] construct vma: [0x100e8, 0x126a0) with flags 0xe for mm 0xffffffffe003ed0b0
[vm.c,122,create_mapping] create_mapping: [0x10000, 0x11000) -> [0x80352000, 0x80353000) with perm 0x5b
[vm.c,122,create_mapping] create_mapping: [0x11000, 0x12000) -> [0x80359000, 0x8035a000) with perm 0x5b
[vm.c,122,create_mapping] create_mapping: [0x12000, 0x13000) -> [0x803ec000, 0x803ed000) with perm 0xdb
[vm.c,287,do_mmap] construct vma: [0x3ffffff00, 0x400000000) with flags 0x7 for mm 0xffffffffe003ed0b0
[vm.c,122,create_mapping] create_mapping: [0x3ffffff00, 0x400000000) -> [0x803eb000, 0x803ec000) with perm
0xd3
[proc.c,235,fork] fork: pid = 4
[trap.c,124,trap_handler] Store/AMO page fault at 0x3fffffffd8
[trap.c,215,do_page_fault] handle COW for process 1 at address 0x3fffffffd8
[trap.c,163,trap_handler] Trap occurs at sepc = 0x1013c
[trap.c,124,trap_handler] Store/AMO page fault at 0x122b0
[trap.c,215,do_page_fault] handle COW for process 1 at address 0x122b0
[trap.c,163,trap_handler] Trap occurs at sepc = 0x10258
[U] pid: 1 is running! global_variable: 2
switch to [PID = 2 PRIORITY = 7 COUNTER = 7]
[trap.c,124,trap_handler] Store/AMO page fault at 0x3fffffffd8
[trap.c,215,do_page_fault] handle COW for process 2 at address 0x3fffffffd8
[trap.c,163,trap_handler] Trap occurs at sepc = 0x1013c
[vm.c,287,do_mmap] construct vma: [0x3ffffff00, 0x400000000) with flags 0x7 for mm 0xffffffffe004370b0
[vm.c,122,create_mapping] create_mapping: [0x3ffffff00, 0x400000000) -> [0x80355000, 0x80356000) with perm
0xd3
[vm.c,287,do_mmap] construct vma: [0x100e8, 0x126a0) with flags 0xe for mm 0xffffffffe004370b0
[vm.c,122,create_mapping] create_mapping: [0x10000, 0x11000) -> [0x80352000, 0x80353000) with perm 0x5b
[vm.c,122,create_mapping] create_mapping: [0x11000, 0x12000) -> [0x80359000, 0x8035a000) with perm 0x5b
[vm.c,122,create_mapping] create_mapping: [0x12000, 0x13000) -> [0x80358000, 0x80359000) with perm 0xdb
[proc.c,235,fork] fork: pid = 5
[trap.c,124,trap_handler] Store/AMO page fault at 0x3fffffffd8
[trap.c,215,do_page_fault] handle COW for process 2 at address 0x3fffffffd8
[trap.c,163,trap_handler] Trap occurs at sepc = 0x1013c
[trap.c,124,trap_handler] Store/AMO page fault at 0x122b0
[trap.c,215,do_page_fault] handle COW for process 2 at address 0x122b0
[trap.c,163,trap_handler] Trap occurs at sepc = 0x1021c
[U] pid: 2 is running! global_variable: 1
[vm.c,287,do_mmap] construct vma: [0x3ffffff00, 0x400000000) with flags 0x7 for mm 0xffffffffe004810b0
[vm.c,122,create_mapping] create_mapping: [0x3ffffff00, 0x400000000) -> [0x8047f000, 0x80480000) with perm
0xd3
[vm.c,287,do_mmap] construct vma: [0x100e8, 0x126a0) with flags 0xe for mm 0xffffffffe004810b0
[vm.c,122,create_mapping] create_mapping: [0x10000, 0x11000) -> [0x80352000, 0x80353000) with perm 0x5b
[vm.c,122,create_mapping] create_mapping: [0x11000, 0x12000) -> [0x80359000, 0x8035a000) with perm 0x5b
[vm.c,122,create_mapping] create_mapping: [0x12000, 0x13000) -> [0x80480000, 0x80481000) with perm 0xdb
[proc.c,235,fork] fork: pid = 6
[trap.c,124,trap_handler] Store/AMO page fault at 0x3fffffffd8
[trap.c,215,do_page_fault] handle COW for process 2 at address 0x3fffffffd8
[trap.c,163,trap_handler] Trap occurs at sepc = 0x1013c
[trap.c,124,trap_handler] Store/AMO page fault at 0x122b0
[trap.c,215,do_page_fault] handle COW for process 2 at address 0x122b0
[trap.c,163,trap_handler] Trap occurs at sepc = 0x10258
[U] pid: 2 is running! global_variable: 2
switch to [PID = 3 PRIORITY = 7 COUNTER = 7]
[trap.c,124,trap_handler] Store/AMO page fault at 0x3fffffffd8

```

图 6: make run TEST=FORK3 proc 1-2

```

switch to [PID = 3 PRIORITY = 7 COUNTER = 7]
[trap.c,124,trap_handler] Store/AMO page fault at 0x3fffffff8
[trap.c,215,do_page_fault] handle COW for process 3 at address 0x3fffffff8
[trap.c,163,trap_handler] Trap occurs at sepc = 0x1013c
[trap.c,124,trap_handler] Store/AMO page fault at 0x122b0
[trap.c,215,do_page_fault] handle COW for process 3 at address 0x122b0
[trap.c,163,trap_handler] Trap occurs at sepc = 0x1021c
[U] pid: 3 is running! global_variable: 1
[vm.c,287,do_mmap] construct vma: [0x3fffffff000, 0x4000000000] with flags 0x7 for mm 0xffffffff0004cc0b0
[vm.c,122,create_mapping] create_mapping: [0x3fffffff000, 0x4000000000] -> [0x803a2000, 0x803a3000] with perm
0xd3
[vm.c,287,do_mmap] construct vma: [0x100e8, 0x126a0] with flags 0xe for mm 0xffffffff0004cc0b0
[vm.c,122,create_mapping] create_mapping: [0x10000, 0x11000] -> [0x80352000, 0x80353000] with perm 0x5b
[vm.c,122,create_mapping] create_mapping: [0x11000, 0x12000] -> [0x80359000, 0x8035a000] with perm 0x5b
[vm.c,122,create_mapping] create_mapping: [0x12000, 0x13000] -> [0x804cb000, 0x804cc000] with perm 0xdb
[proc.c,235,fork] fork: pid = 7
[trap.c,124,trap_handler] Store/AMO page fault at 0x3fffffff8
[trap.c,215,do_page_fault] handle COW for process 3 at address 0x3fffffff8
[trap.c,163,trap_handler] Trap occurs at sepc = 0x1013c
[trap.c,124,trap_handler] Store/AMO page fault at 0x122b0
[trap.c,215,do_page_fault] handle COW for process 3 at address 0x122b0
[trap.c,163,trap_handler] Trap occurs at sepc = 0x10258
[U] pid: 3 is running! global_variable: 2
switch to [PID = 4 PRIORITY = 7 COUNTER = 7]
[trap.c,124,trap_handler] Store/AMO page fault at 0x3fffffff8
[trap.c,215,do_page_fault] handle COW for process 4 at address 0x3fffffff8
[trap.c,163,trap_handler] Trap occurs at sepc = 0x1013c
[trap.c,124,trap_handler] Store/AMO page fault at 0x122b0
[trap.c,215,do_page_fault] handle COW for process 4 at address 0x122b0
[trap.c,163,trap_handler] Trap occurs at sepc = 0x10258
[U] pid: 4 is running! global_variable: 2
switch to [PID = 5 PRIORITY = 7 COUNTER = 7]
[trap.c,124,trap_handler] Store/AMO page fault at 0x3fffffff8
[trap.c,215,do_page_fault] handle COW for process 5 at address 0x3fffffff8
[trap.c,163,trap_handler] Trap occurs at sepc = 0x1013c
[trap.c,124,trap_handler] Store/AMO page fault at 0x122b0
[trap.c,215,do_page_fault] handle COW for process 5 at address 0x122b0
[trap.c,163,trap_handler] Trap occurs at sepc = 0x1021c
[U] pid: 5 is running! global_variable: 1
[vm.c,287,do_mmap] construct vma: [0x100e8, 0x126a0] with flags 0xe for mm 0xffffffff0005160b0
[vm.c,122,create_mapping] create_mapping: [0x10000, 0x11000] -> [0x80352000, 0x80353000] with perm 0x5b
[vm.c,122,create_mapping] create_mapping: [0x11000, 0x12000] -> [0x80359000, 0x8035a000] with perm 0x5b
[vm.c,122,create_mapping] create_mapping: [0x12000, 0x13000] -> [0x80358000, 0x80359000] with perm 0xdb
[vm.c,287,do_mmap] construct vma: [0x3fffffff000, 0x4000000000] with flags 0x7 for mm 0xffffffff0005160b0
[vm.c,122,create_mapping] create_mapping: [0x3fffffff000, 0x4000000000] -> [0x80355000, 0x80356000] with perm
0xd3
[proc.c,235,fork] fork: pid = 8
[trap.c,124,trap_handler] Store/AMO page fault at 0x3fffffff8
[trap.c,215,do_page_fault] handle COW for process 5 at address 0x3fffffff8
[trap.c,163,trap_handler] Trap occurs at sepc = 0x1013c
[trap.c,124,trap_handler] Store/AMO page fault at 0x122b0
[trap.c,215,do_page_fault] handle COW for process 5 at address 0x122b0
[trap.c,163,trap_handler] Trap occurs at sepc = 0x10258
[U] pid: 5 is running! global_variable: 2
switch to [PID = 6 PRIORITY = 7 COUNTER = 7]
[trap.c,124,trap_handler] Store/AMO page fault at 0x3fffffff8
[trap.c,215,do_page_fault] handle COW for process 6 at address 0x3fffffff8
[trap.c,163,trap_handler] Trap occurs at sepc = 0x1013c
[trap.c,124,trap_handler] Store/AMO page fault at 0x122b0
[trap.c,215,do_page_fault] handle COW for process 6 at address 0x122b0
[trap.c,163,trap_handler] Trap occurs at sepc = 0x10258
[U] pid: 6 is running! global_variable: 2
switch to [PID = 7 PRIORITY = 7 COUNTER = 7]
[trap.c,124,trap_handler] Store/AMO page fault at 0x3fffffff8
[trap.c,215,do_page_fault] handle COW for process 7 at address 0x3fffffff8
[trap.c,163,trap_handler] Trap occurs at sepc = 0x1013c
[trap.c,124,trap_handler] Store/AMO page fault at 0x122b0
[trap.c,215,do_page_fault] handle COW for process 7 at address 0x122b0
[trap.c,163,trap_handler] Trap occurs at sepc = 0x10258
[U] pid: 7 is running! global_variable: 2
switch to [PID = 8 PRIORITY = 7 COUNTER = 7]
[trap.c,124,trap_handler] Store/AMO page fault at 0x3fffffff8
[trap.c,215,do_page_fault] handle COW for process 8 at address 0x3fffffff8
[trap.c,163,trap_handler] Trap occurs at sepc = 0x1013c
[trap.c,124,trap_handler] Store/AMO page fault at 0x122b0
[trap.c,215,do_page_fault] handle COW for process 8 at address 0x122b0
[trap.c,163,trap_handler] Trap occurs at sepc = 0x10258
[U] pid: 8 is running! global_variable: 2
switch to [PID = 1 PRIORITY = 7 COUNTER = 7]

```

图 7: make run TEST=FORK3 proc 3-8

这与上面我们推测的进程树结果是一致的.